



AUGUST 9-10, 2023

BRIEFINGS

Dive into Apple UserFS (Userspace Filesystem)

pattern-f (@pattern_F_)
Ant Security Light-Year Lab

About me



- pattern-f (@pattern_F_) on Twitter
- Security researcher of Ant Security Light-Year Lab
- Focus on iOS & macOS security
- speaker of Black Hat ASIA & USA 2021

Background

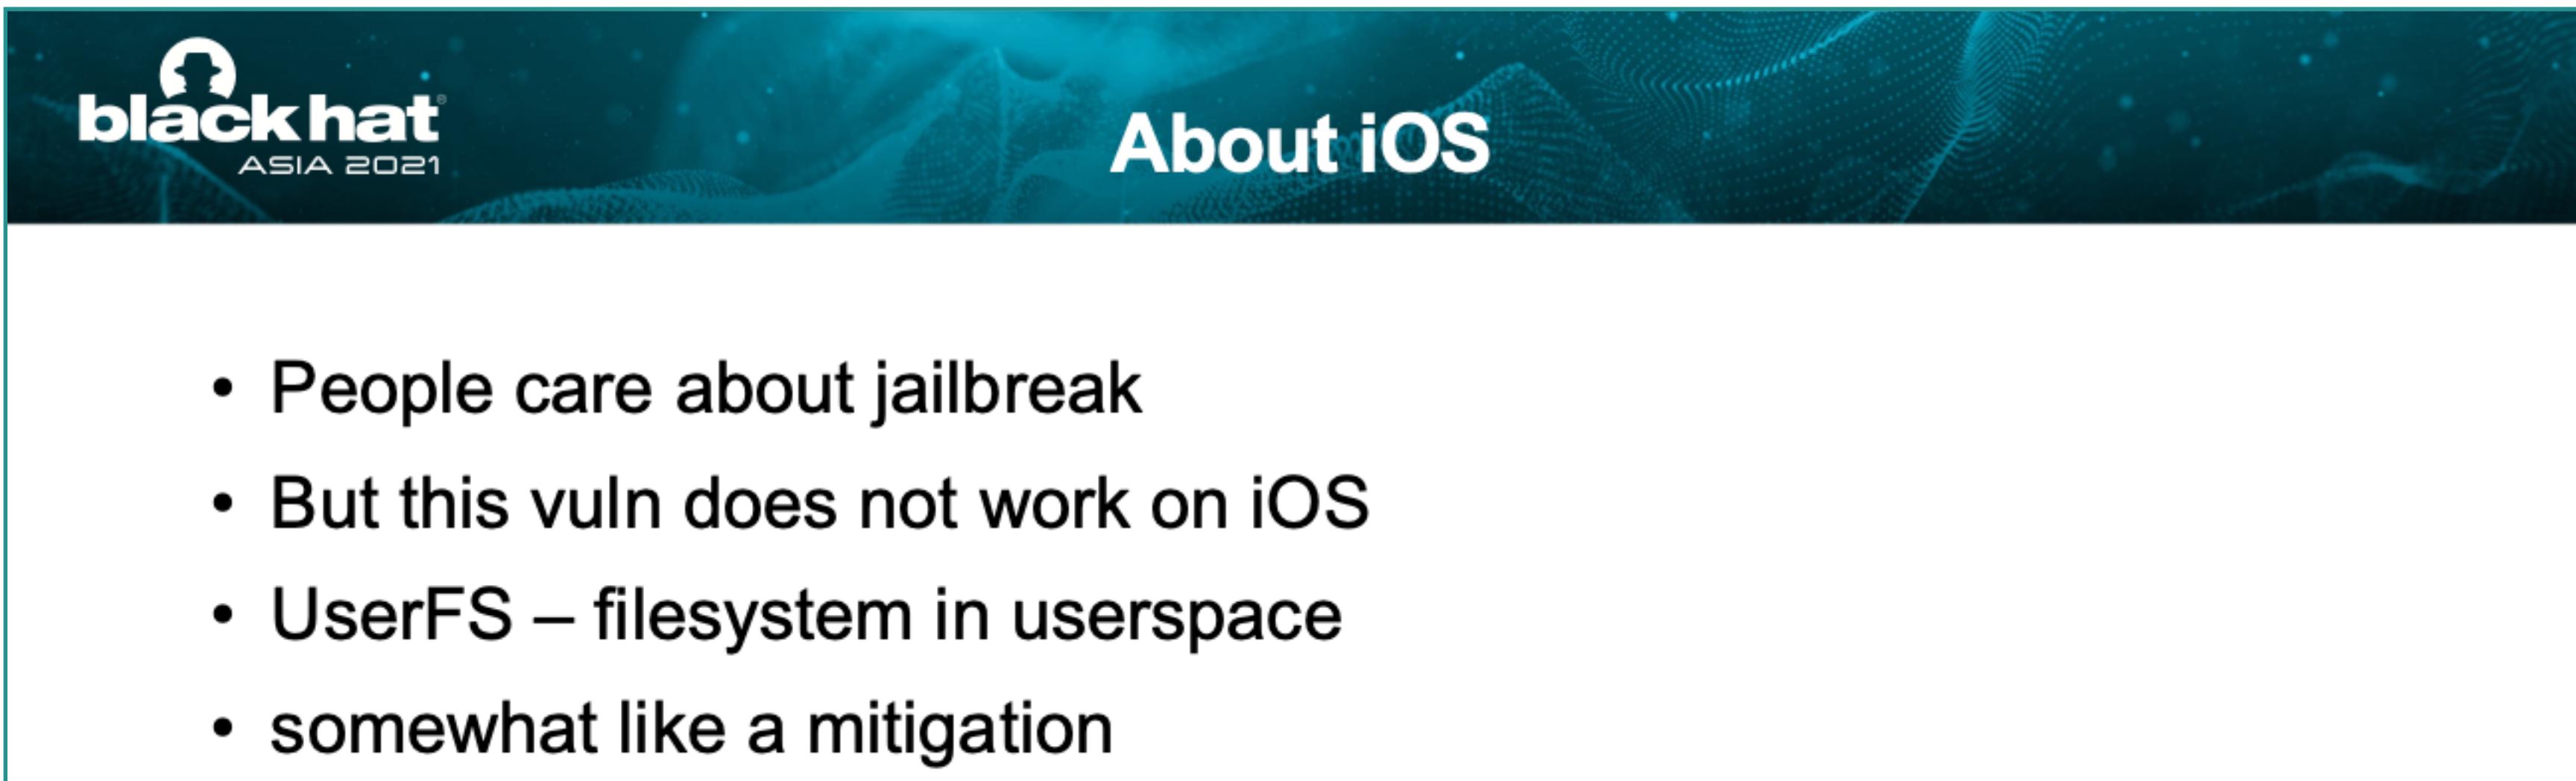
- my previous talk at Black Hat ASIA 2021

The Price of Compatibility: Defeating macOS Kernel Using Extended File Attributes
Speaker: Zuozhi Fan

- Shared a vulnerability named xattr-oob-swap (CVE-2020-27904) [\[link\]](#)
- It's a filesystem bug, can get tfp0 on macOS 10.15.x and below.
- iOS & macOS share code base, it should work on iOS too, but...

Background

- It doesn't crash iOS. Maybe UserFS stops it.
- I don't know the mechanism of UserFS, so I'm not sure.



The slide has a dark teal header with the Black Hat Asia 2021 logo on the left and the title "About iOS" in white on the right. The main content area is white with a teal border.

- People care about jailbreak
- But this vuln does not work on iOS
- UserFS – filesystem in userspace
- somewhat like a mitigation

Background



- This vulnerability doesn't affect iOS, really?
- Add it to TODO list



- This vulnerability doesn't affect iOS, really?
- Add it to ~~TODO~~ list

File System

Available for: iPhone 6s (all models), iPhone 7 (all models), iPhone SE (1st generation), iPad Pro (all models), iPad Air 2 and later, iPad 5th generation and later, iPad mini 4 and later, and iPod touch (7th generation)

Impact: An app may be able to break out of its sandbox

Description: This issue was addressed with improved checks.

CVE-2022-42861: pattern-f (@pattern_F_) of Ant Security Light-Year Lab

What is UserFS

- We know that



Files (Apple)

[Article](#) [Talk](#)

From Wikipedia, the free encyclopedia

Files is a [file management app](#) developed by [Apple Inc.](#) for devices that run [iOS 11](#) and later releases of [iOS](#) and devices that run [iPadOS](#). Discovered as a placeholder title in the [App Store](#)

- iPhone User Guide said: An external storage device must have only a single data partition, and it must be formatted as **APFS**, APFS (encrypted), macOS Extended (**HFS+**), **exFAT** (FAT64), **FAT32**, or **FAT**.

What is UserFS

- iOS supports various file system
- But when analyzing the iOS kernel cache, I found only apfs & hfs support were present.
- No msdos.kext (fat32, exfat) at all

	com.apple.filesystems.hfs.kext:HEADER
	com.apple.filesystems.hfs.kext:__const
	com.apple.filesystems.hfs.kext:__cstring
	com.apple.filesystems.apfs:HEADER
	com.apple.filesystems.apfs:__const
	com.apple.filesystems.apfs:__cstring
	com.apple.filesystems.lifs:HEADER
	com.apple.filesystems.lifs:__os_log
	com.apple.filesystems.lifs:__cstring

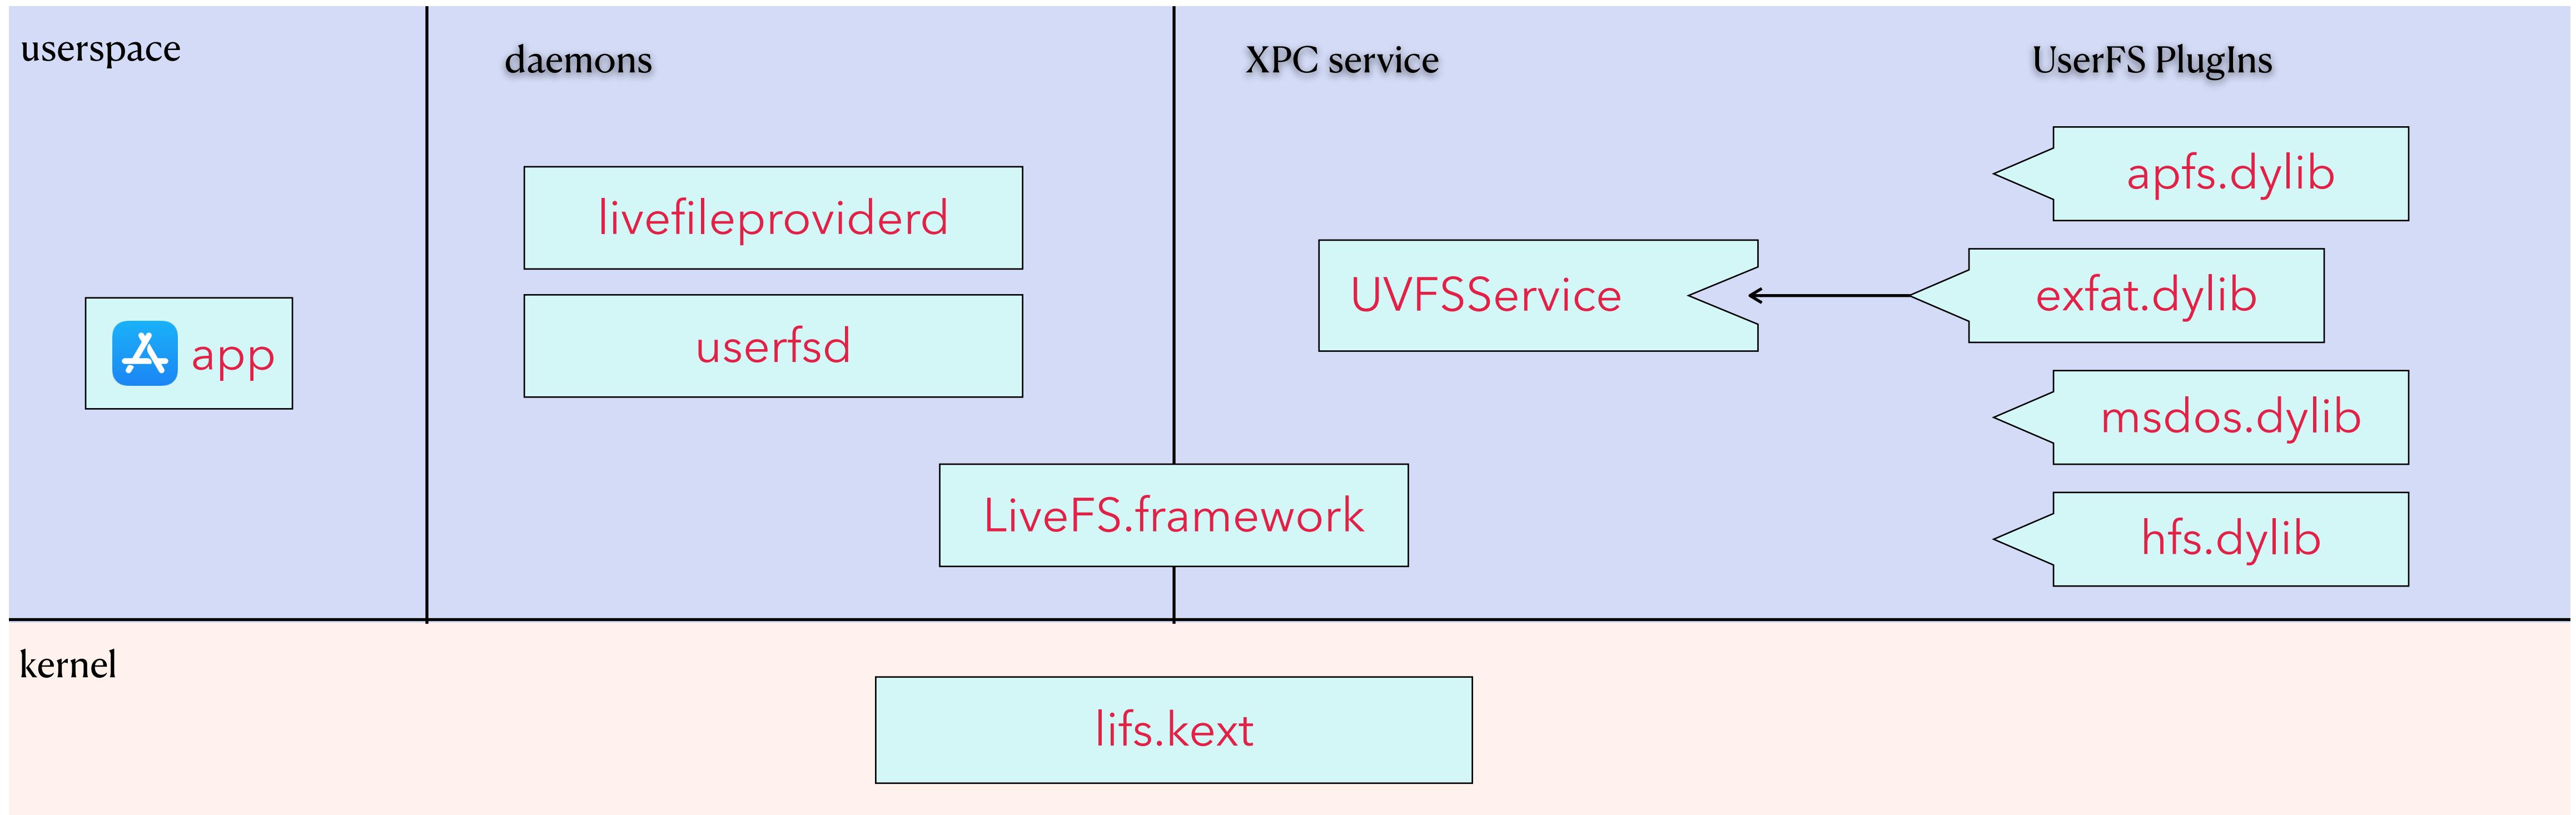
What is UserFS

- Instead, apple implement them in user space, called UserFS plugins.
- But the plugins are not all about UserFS.

```
/System/Library/PrivateFrameworks/UserFS.framework/PlugIns/
└── liblivefiles.plugin.dummy.dylib
└── livefiles_apfs.dylib
└── livefiles_cs.dylib
└── livefiles_exfat.dylib
└── livefiles_hfs.dylib
└── livefiles_msdos.dylib
```

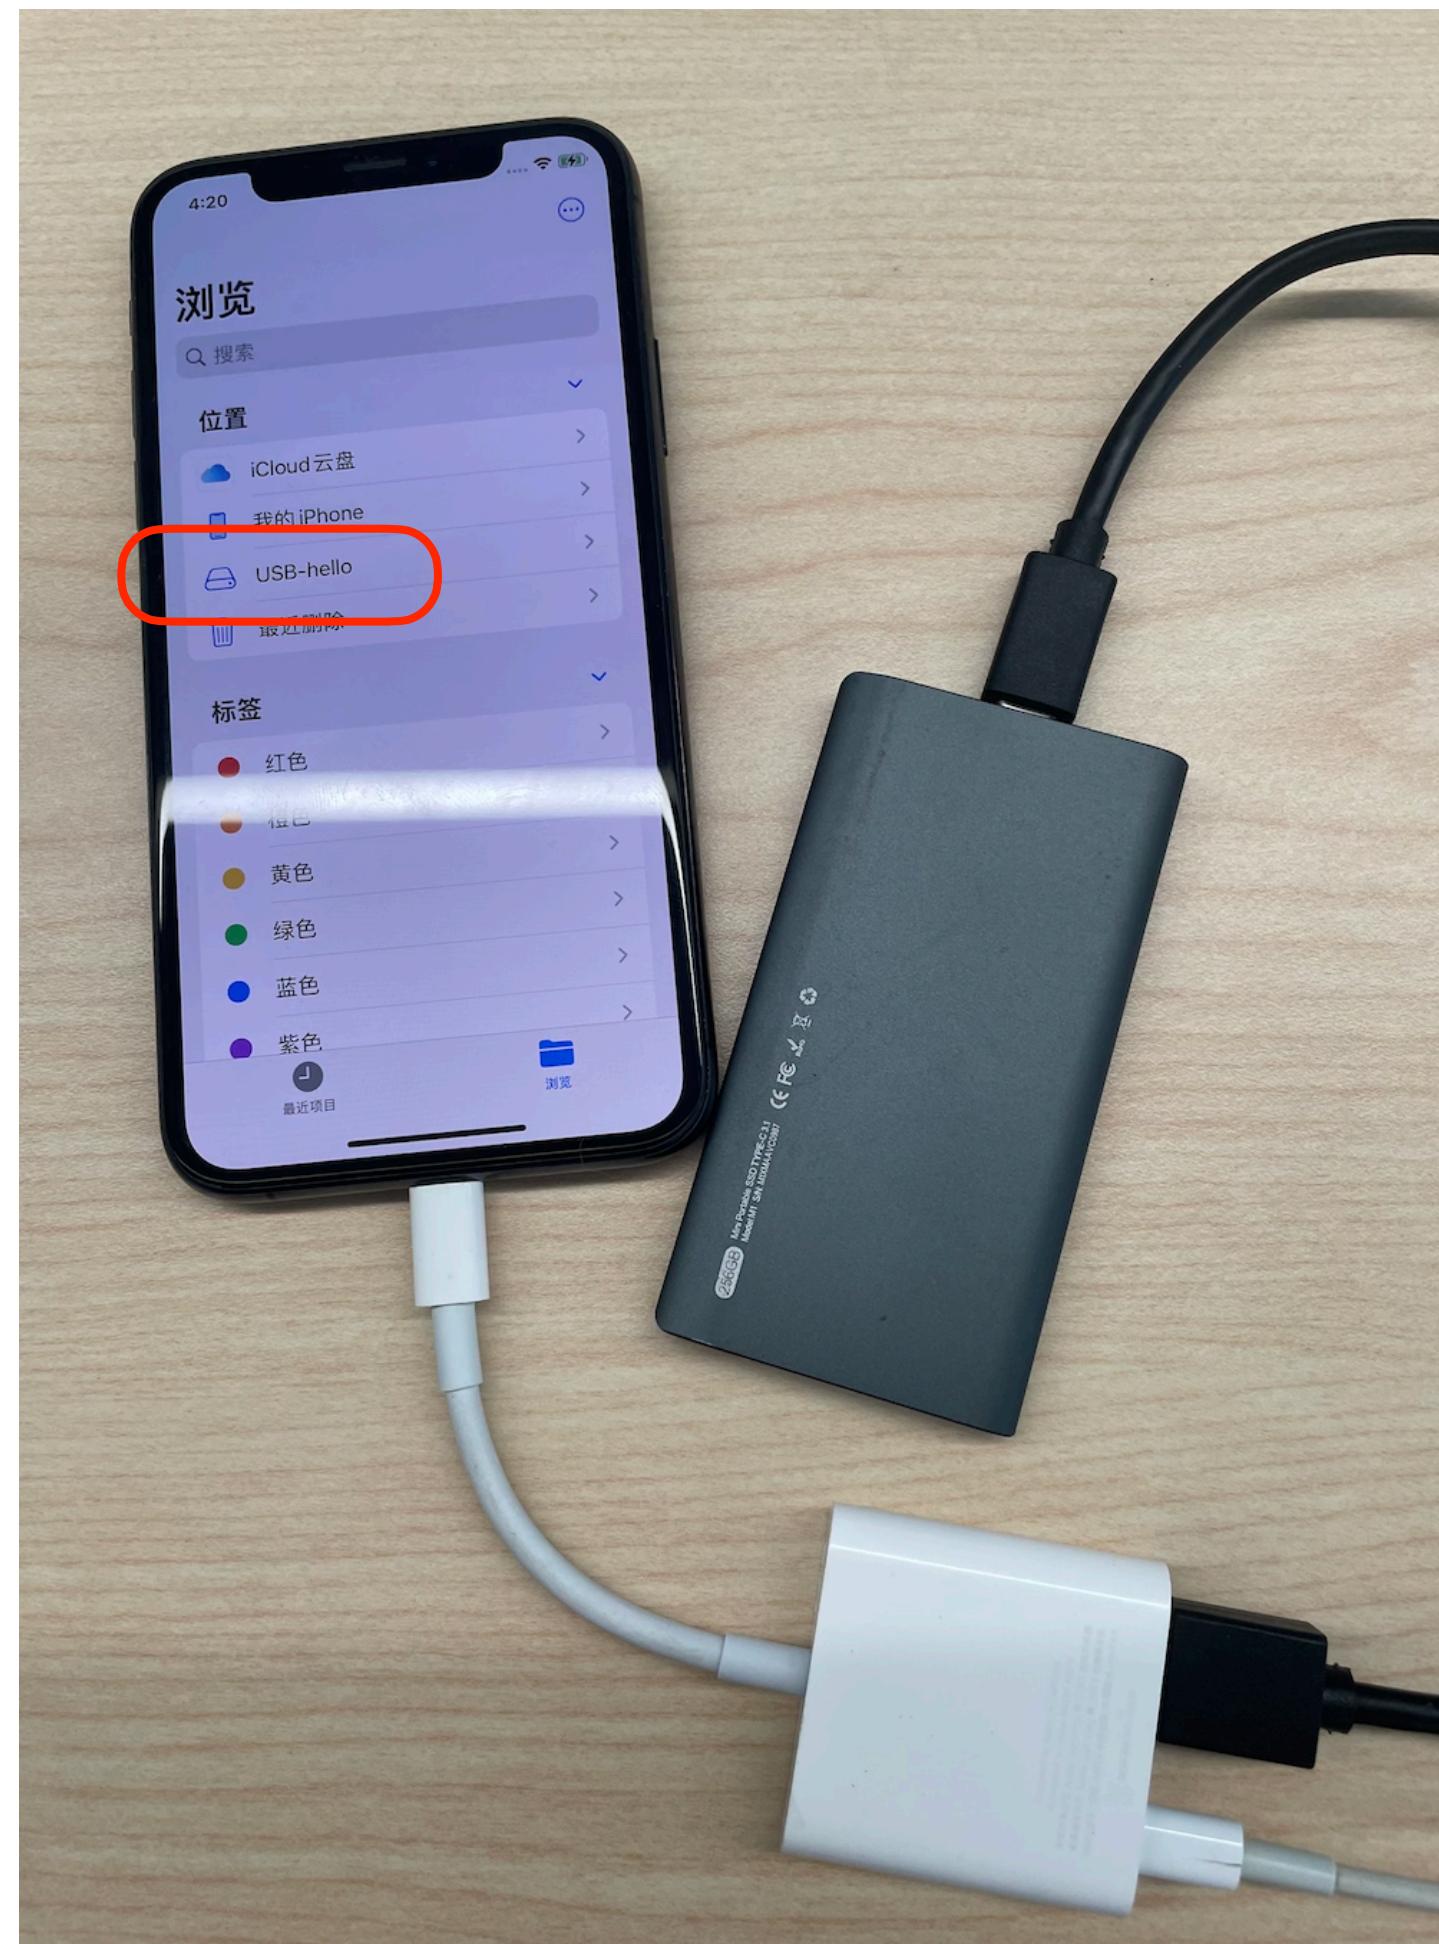
What is UserFS

- UserFS is filesystem implemented in user space, with these components.
 - kext, daemon, XPC service, PlugIn, and framework



What is UserFS

- UserFS only functions with external storage devices, so let's start with connecting an external hard drive to iPhone.



When connecting a USB drive to iPhone

- launchd registers a LaunchEvent “usb.device.attached” for userfsd

```
<key>/System/Library/LaunchDaemons/com.apple.filesystems.userfsd.plist</key>
<dict>
  <key>LaunchEvents</key>
  <dict>
    <key>com.apple.iokit.matching</key>
    <dict>
      <key>com.apple.userfsd.usb.device.attached</key>
      <dict>
        <key>IOProviderClass</key>
        <string>IOBlockStorageDevice</string>
        <key>IOMatchLaunchStream</key>
        <true/>
        <key>IOParentMatch</key>
        <dict>
          <key>IOProviderClass</key>
          <string>IOUSBHostDevice</string>
          <key>IOPropertyMatch</key>
          <dict>
            <key>removable</key>
            <true/>
```

When connecting a USB drive to iPhone

- userfsd daemon starts and handles this IOKit notification
- Make a xpc call to UVFSService

```
void handleIOKitNotifications(void)
{
    service_info = IOServiceMatching("IOMedia");
    IOServiceAddMatchingNotification(gNotifyPort, "IOServiceFirstMatch", service_info,
        IOMediaMatchingCallback, 0, &iterator);
}

void IOMediaMatchingCallback(void *refcon, io_iterator_t iterator)
{
    diskName = CFDictionaryGetValue(properties, CFSTR("BSD Name"));
    objc_msgSend(externalVolumeLiveFSService, "LiveMountAddDisk:reply:", diskName, reply_b);
}

void _[externalVolumeLiveFSServiceDelegate LiveMountService:addDisk:reply:] ✘ Expected identifier
{
    -[NSXPConnection initWithServiceName:CFSTR("com.apple.filesystems.lifs.userfsd.UVFSService")];
    objc_msgSend(remoteObjectProxy, "createVolumesForTheDevice:how:withReply:", v30, v43, &v54);
}
```

/dev/disk3

When connecting a USB drive to iPhone

- Enumerate UserFS plugins and determine filesystem type of the disk.
- xc call again, to livefileproviderd.

```
-[UVFSService createVolumesForTheDevice:how:withReply:]  
{  
    -[UVFSService CreateVolumesForTheDevice:how:] {  
        rawDevice = -[LiveFSRawDevice initDeviceWithName:andError:]  
        volumes = [rawDevice getVolumesFromDevice] {  
            _fsPlugin = -[UVFSPlugin loadFileSystemDyLibPassingParameterDict:forShimPlugin:] {  
                dyLibHandle = dlopen("/PrivateFrameworks/UserFS.framework/PlugIns/livefiles_%s.dylib")  
                dlsym(dyLibHandle, "livefiles_shim_plugin_init")(&self->_FSOps)  
            }  
            FSOps = -[UVFSPlugin FSOps](self->_fsPlugin, "FSOps");  
            FSOps->fsops_taste(self->_deviceFD);  
            FSOps->fsops_scavols(self->_deviceFD)  
        }  
        mountClient = +[LiveFSMountClient newClientForProvider:] {  
            self->conn = connectionForServiceURL("machp://com.apple.filesystems.livefileproviderd");  
        }  
        [mountClient mountVolume:displayName:provider:domainError:on:how:] {  
            [proxyObject LiveMounterMountVolume:displayName:provider:domainError:on:how:reply:]  
        }  
    }  
}
```

When connecting a USB drive to iPhone

- familiar command “/sbin/mount”, mount_lifs this time.

```
-[lifeFilesFPNFSDMounter LiveMounterMountVolume:displayName:provider:domainError:on:how:reply:]  
{  
    -[lifeFilesFPNFSDMounter LiveMounterReallyMountVolume:displayName:provider:domainError:on:how:reply:]  
    ...  
    -[mountEntry mount:] {  
        posix_spawn(&pid, "/sbin/mount_lifs", actions, attr, { "/sbin/mount_lifs", ... }, environ)  
    }  
}  
  
/sbin/mount_lifs  
{  
    mount("lifs", dir, flags, data) {  
        lifs.kext`lifs_mount(...)  
    }  
}
```

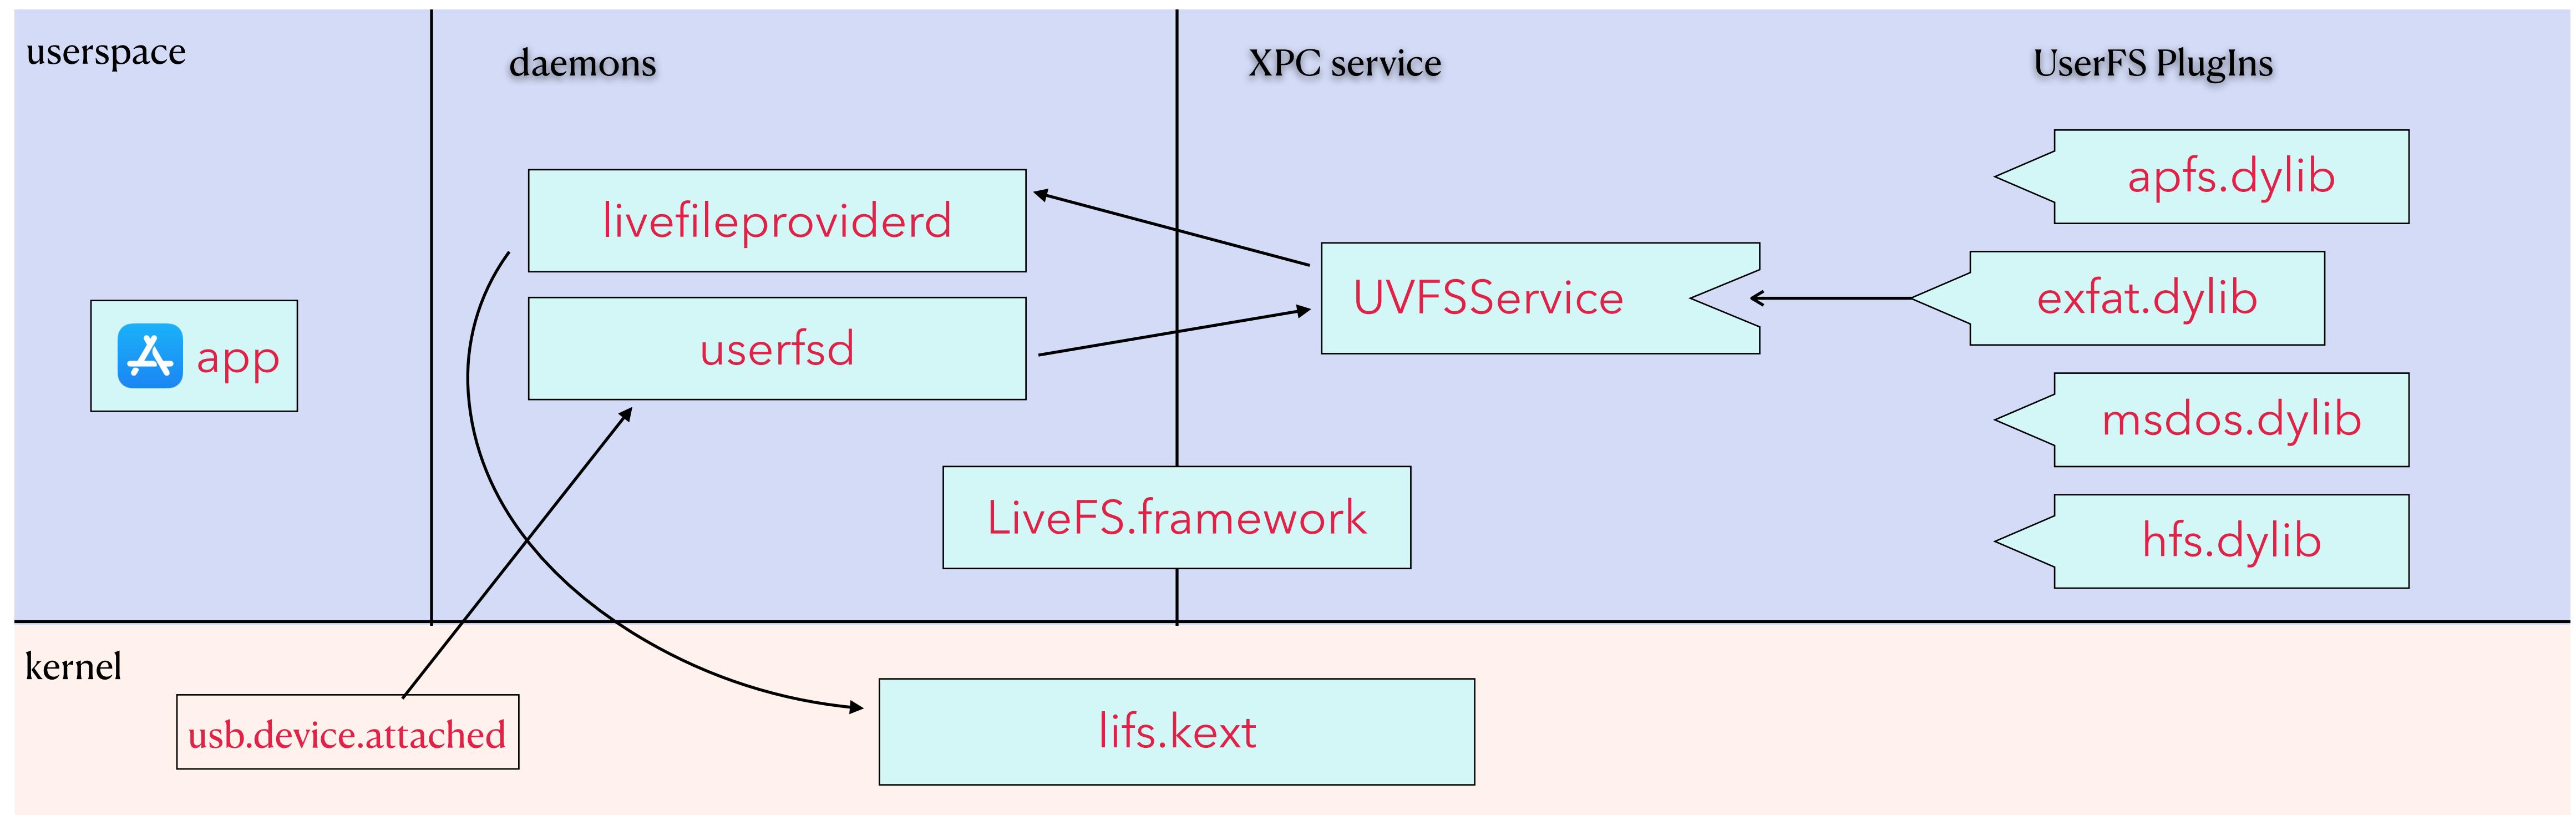
When connecting a USB drive to iPhone

- Finally, we can access files on this exFAT formatted USB hard drive.
- /var/mobile/Library/LiveFiles/ is a dedicated directory.

```
bash.dylib-5.2# /sbin/mount  
lifs://livefileproviderd@291/LiveFiles on /private/var/mobile/Library/LiveFiles (lifs, nodev, noexec, nosuid)  
exfat://disk3s1/USB-hello on /private/var/mobile/Library/LiveFiles/com.apple.filesystems.userfsd/USB-hello (exfat, nodev, noexec, nosuid)  
bash.dylib-5.2# ls --color /private/var/mobile/Library/LiveFiles/com.apple.filesystems.userfsd/USB-hello  
1.txt 2.txt 'System Volume Information' autorun.inf eaget.ico temp uaf.txt
```

When connecting a USB drive to iPhone

- Summary: process of mounting a UserFS volume



When accessing files via UserFS

- To read a file: `int fd = open(path, O_RDONLY); read(fd, ...);`
- We know that `SYS_open` will be dispatched to `vnop_open`

```
(lldb) bt
* thread #2, stop reason = breakpoint 2.1
* frame #0: 0xffffffff009e084b4 ios15-kc.out`hfs_vnop_open
  frame #1: 0xffffffff007fdf4b0 ios15-kc.out`vn_open_auth + 1884
  frame #2: 0xffffffff007fc822c ios15-kc.out`open1 + 244
  frame #3: 0xffffffff007fc8e70 ios15-kc.out`open_nocancel + 268
  frame #4: 0xffffffff008423e0c ios15-kc.out`unix_syscall + 756
```

When accessing files via UserFS

- Accessing files via UserFS is same, except that SYS_open will be dispatched to lifs.kext`lifs_vnop_open

```
int64 __fastcall lifs_vnop_open(vnop_open_args *a)
{
    struct vnode *a_vp; // x19
    char *fsnode; // x20
    __int64 err; // x21
    int a_mode; // w26
    int v6; // w23
    uintptr_t fsmount; // [xsp+8h] [xbp-48h] BYREF

    a_vp = a->a_vp;
    fsnode = (char *)vnode_fsnodes(a_vp);
    fsmount = 0LL;
    err = get_lifs_mount_from_node((lifsnode *)fsnode, (lifsmount **)&fsmount);
    if ( !(_DWORD)err )
    {
        if ( (kdebug_enable & 0xFFFFFFFF7) != 0 )
            kernel_debug_filtered(0x3140089u, fsmount, (uintptr_t)a_vp, (uintptr_t)fsnode, a
        a_mode = a->a_mode;
        v6 = a_mode & 3;
        err = lifs_open_request(fsmount, (lifsnode *)fsnode, a_mode & 3);
    }
}
```

When accessing files via UserFS

- vnop_open will build a mach message and send it to a server.

```
int lifs_open_request(lifsmount *li_mount, void *fsnode, void *a3)
{
    get_lifs_port(&svr_port);

    req.request_id = OSAddAtomic64(1LL, &lifs_request_id);
    // ... copy params to request

    lifs_add_req(&req);

    ret = lifs_open_send(svr_port, req.request_id, fsnode, a3); {
        struct { mach_msg_header_t msgHdr; } openRequest;
        openRequest.msgHdr.msgh_id = 0x2A5;
        mach_msg_send_from_kernel_proper(&openRequest.msgHdr, sizeof(openRequest));
    }
    if ( ret == 0 ) {
        lifs_wait_req_completion(&req);
        ret = req.retcode_2C;
    }

    lifs_remove_req(&req);
    return ret;
}
```

When accessing files via UserFS

- Find handler for the open request.
- livefileproviderd registers a notification port. That's the server port.

```
livefileproviderd`main()
{
    mach_port_allocate(mach_task_self, 1, &svr_port);

    li_UserClient = +[LiveFSUserClient defaultClient];
    objc_msgSend(li_UserClient, "setMainMachPort:forDomain:", svr_port) {
        IOServiceOpen(IOServiceMatching("com_apple_filesystems_lifs"), mach_task_self, 0, &self->ourPort);
        IOConnectSetNotificationPort(self->ourPort, 0, svr_port, 0LL);
    }
}

AppleLIFSUserClient::registerNotificationPort(AppleLIFSUserClient *this, ipc_port_t svr_port)
{
    return lifs_set_machport(svr_port, this->clientDomain);
}
```

When accessing files via UserFS

- livefileproviderd starts a mig server.

```
livefileproviderd`main()
{
    source = dispatch_source_create(DISPATCH_SOURCE_TYPE_MACH_RECV, svr_port, ...);
    dispatch_source_set_event_handler(source, ^{
        dispatch_mig_server(source, 2168, mig_message_dispatcher);
    });
}

void *mig_lifs_open_send(void *InHeadP, void *OutHeadP)
{
    lifs_open_send(request_id, fsnode, mode, token) {
        fileHandle = m_resolve_fsnodes(fsnodes, &a2, &mountEntry, &a4, 0);
        id fsObj = [mountEntry fsObjWithErrorHandler:];
        objc_msgSend(fsObj, "LIOpen:withMode:forPID:reply:", fileHandle, mode, pid, ^(int ret){
            int selector = 2;
            lifs_send_reply(request_id, ret, selector, (_int64)&v2, 16LL) {
                objc_msgSend(li_UserClient, "callStructMethod:inStruct:inSize:outStruct:outStructSize:");
            }
        });
    }
}
```

A red box highlights the line `openRequest.msgHdr.msgh_id = 0x2A5;`. A red arrow points from this box to the line `objc_msgSend(fsObj, "LIOpen:withMode:forPID:reply:", fileHandle, mode, pid, ^(int ret){`.

A red box highlights the line `IOConnectCallStructMethod`. A red arrow points from the bottom of the previous red box down to this line.

When accessing files via UserFS

- `-[LiveFSUserClient callStructMethod:inStruct:inSize:outStruct:outStructSize:]`

```
lifs.kext`AppleLIFSUserClient::methodGenericReply(
    AppleLIFSUserClient *a1,
    void *a2,
    IOExternalMethodArguments *a3)
{
    lifs_request_done(*(_QWORD *)a3->structureInput, *((_DWORD *)a3->structureInput + 2), 0, 0, 0, 0);
    return 0;
}
```



When accessing files via UserFS

```
int lifs_open_request(lifsmount *li_mount, void *fsnode, void *a3)
{
    get_lifs_port(&svr_port);

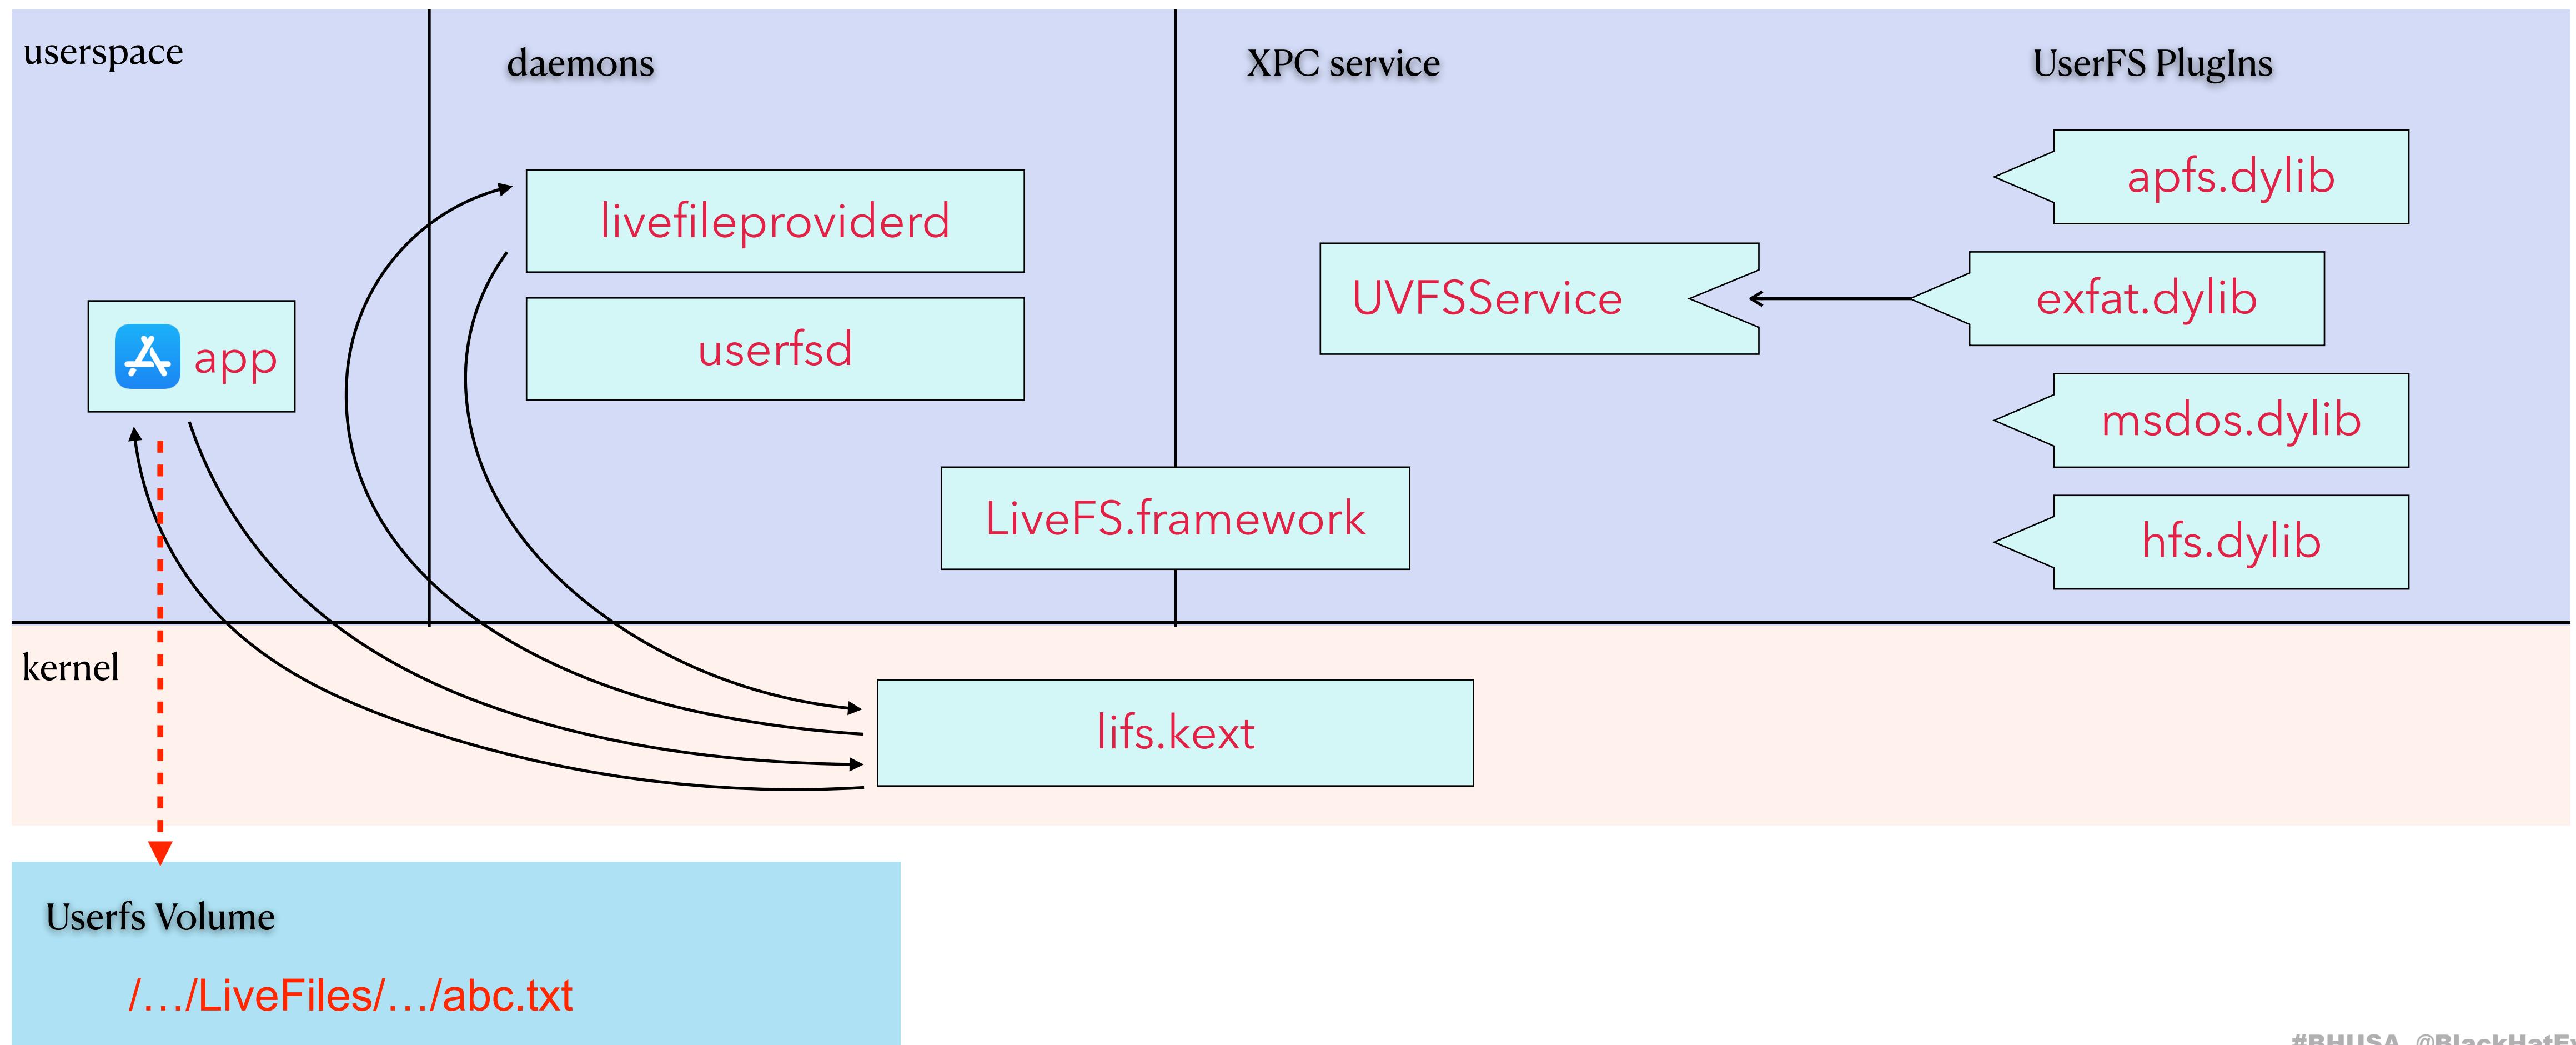
    req.request_id = OSAddAtomic64(1LL, &lifs_request_id);
    // ...

    lifs_add_req(&req);

    ret = lifs_open_send(svr_port, req.request_id, fsnode, a3); {
        struct { mach_msg_header_t msgHdr; } openRequest;
        openRequest.msgHdr.msgh_id = 0x2A5;
        mach_msg_send_from_kernel_proper(&openRequest.msgHdr, sizeof(openRequest));
    }
    if ( ret == 0 ) {
        lifs_request_done(...);
        lifs_wait_req_completion(&req);
        int fd = open(path, O_RDONLY);
        ret = req.retcode_2C;
    }
    lifs_remove_req(&req);
    return ret;
}
```

When accessing files via UserFS

- summarize control flow of “open” syscall



When accessing files via UserFS

- livefileproviderd

```
objc_msgSend(v15, "LIOpen:withMode:forPID:reply:", fileHandle, mode, pid, &v24);
```

- LiveFS.framework

```
void __cdecl -[LiveFSServiceConnection LIOpen:withMode:forPID:reply:](  
    LiveFSServiceConnection *self,  
    SEL a2,  
    id a3,  
    int a4,  
    int a5,  
    id a6)  
{  
    -[NSFileProviderLiveItemImplementation LIOpen:withMode:forPID:reply:](  
        self->mount,  
        "LIOpen:withMode:forPID:reply:",  
        a3, off=8; NSFileProviderLiveItemImplementation *  
        *( QWORD *)&a4,
```

- UVFSService

```
void __cdecl -[liveFSVolume LIOpen:withMode:forPID:reply:](liveFSVolume *self, SEL a2, id
```

When accessing files via UserFS

- UVFSService

```
void -[liveFSVolume LILookup:name:forClient:reply:](  
    liveFSVolume *self, SEL a2, id fileHandle, id nameStr, unsigned client, id reply_b)  
{  
    fileNode = -[liveFSVolume getNodeForFH:fileHandle withError:error];  
    [fileNode lookup:nameStr withResultingNode:&resultNode]; {  
        self->FSOps->fsops_lookup(self->_UVFSNode, [nameStr UTF8String], &uvfsNode);  
        fileNode = [[liveFSNode alloc] initWithVolume:self->volume  
                                              andParent:self  
                                              andName:nameStr  
                                              andUVFSNode:uvfsNode];  
        *resultNode = fileNode;  
    }  
    resultFileHandle = [resultNode getFH];  
    attrData = [resultNode getAttrData];  
    reply_b(retcode, resultFileHandle, attrData);  
}
```

```
void __cdecl -[liveFSVolume LIOpen:withMode:forPID:reply:](liveFSVolume *self, SEL a2, id
```

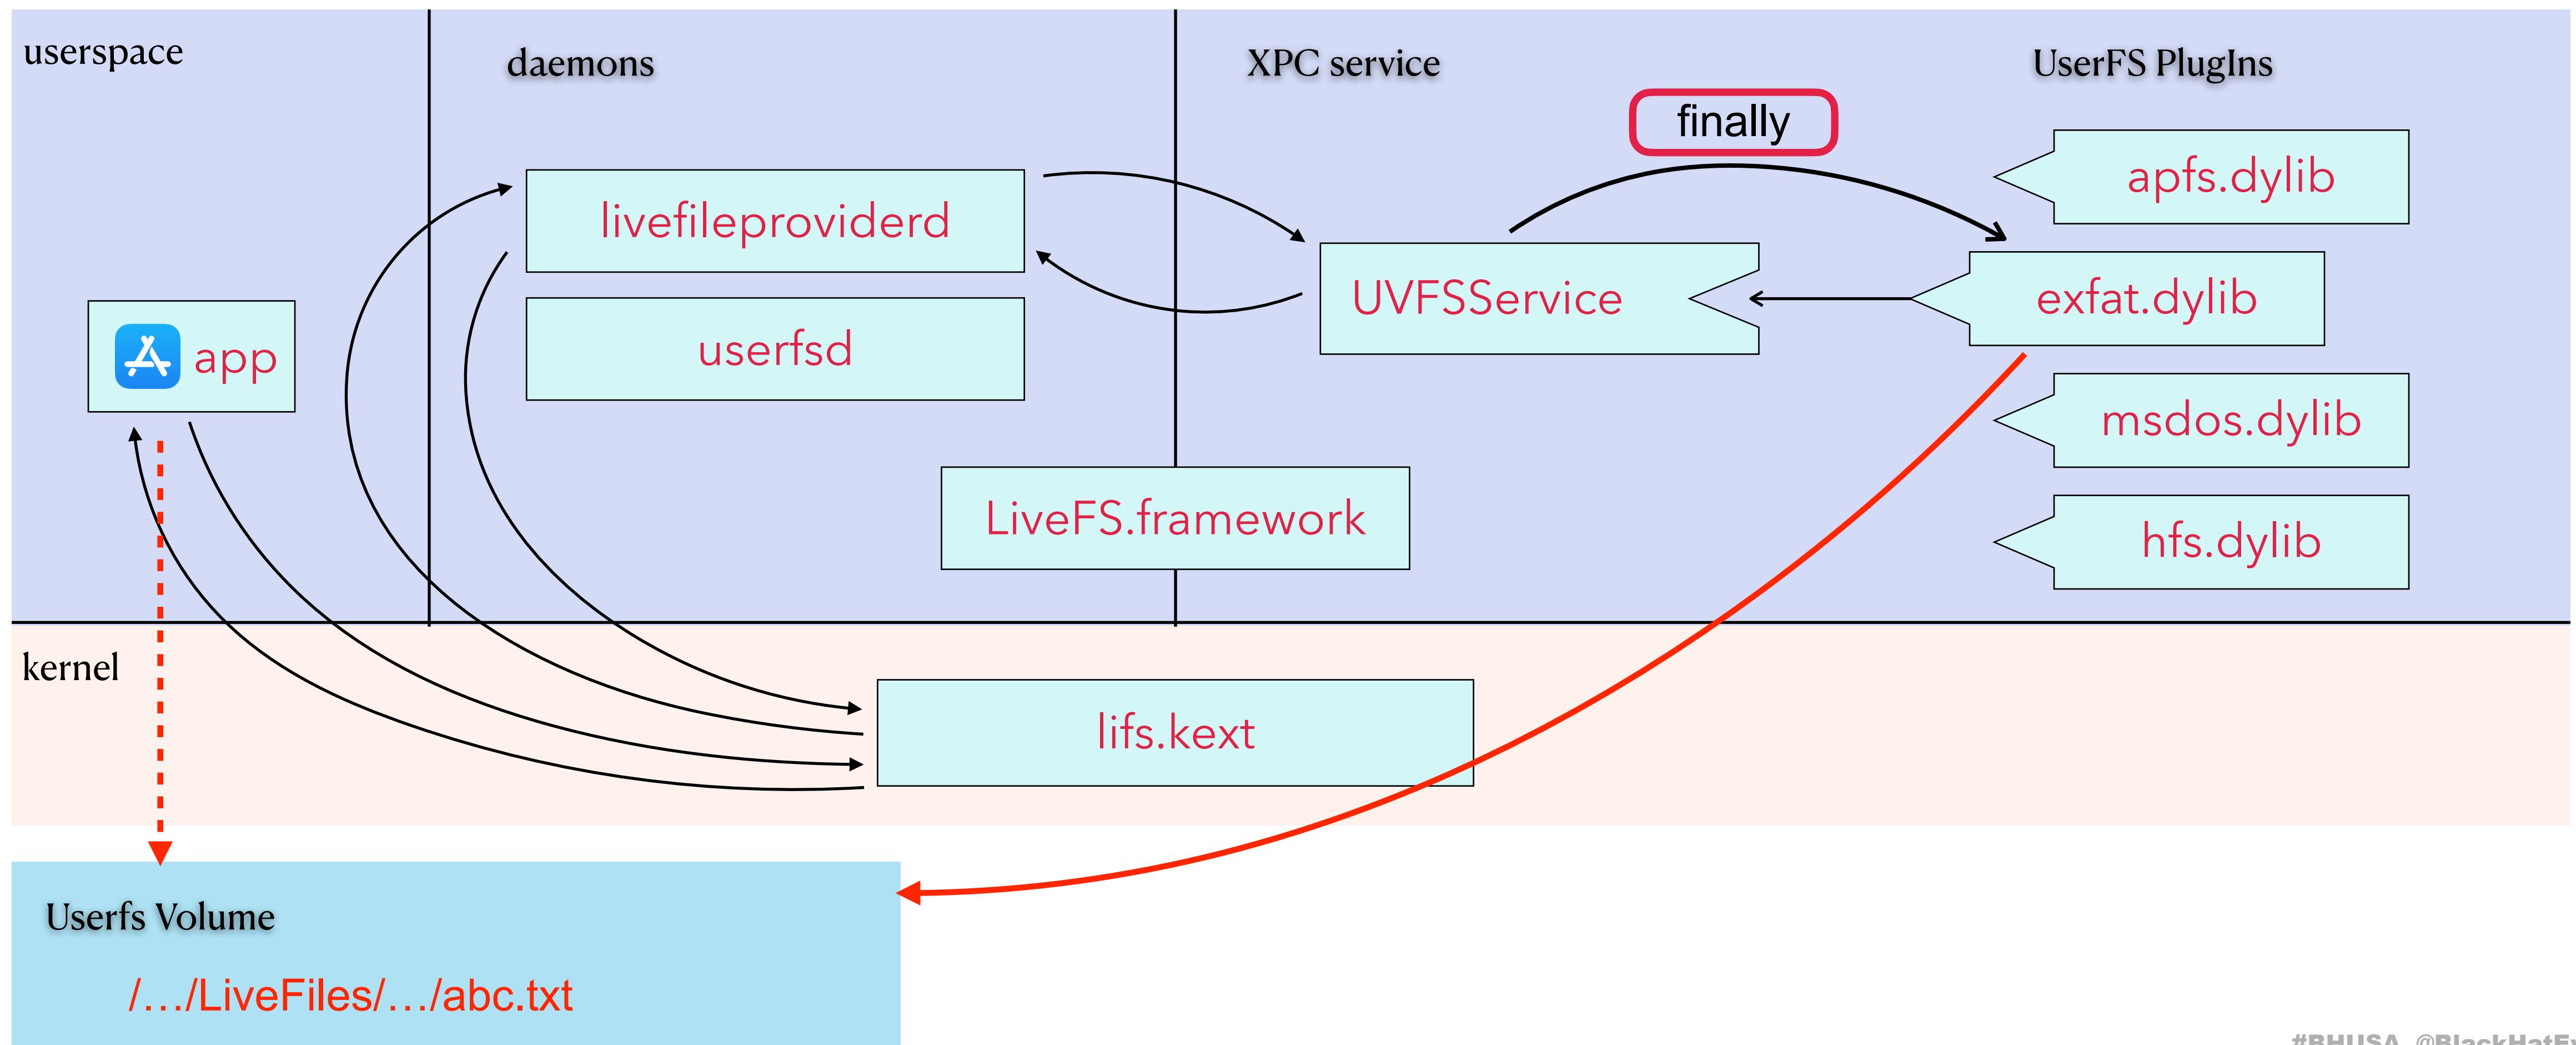
When accessing files via UserFS

- UserFS.framework/PlugIns/**livefiles_exfat.dylib** will do the real job.
- Parse /dev/disk3s1 as a normal file (or disk image), execute exfat file read/write request.
- Functions exported by exfat.dylib:

 <code>_EXFAT_GetAttr</code>
 <code>_EXFAT_GetFSAttr</code>
 <code>_EXFAT_Init</code>
 <code>_EXFAT_Link</code>
 <code>_EXFAT_LoggerInit</code>
 <code>_EXFAT_Lookup</code>
 <code>_EXFAT_MkDir</code>
 <code>_EXFAT_Mount</code>
 <code>_EXFAT_Read</code>
 <code>_EXFAT_ReadDir</code>
 <code>_EXFAT_ReadDirAttr</code>
 <code>_EXFAT_ReadLink</code>
 <code>_EXFAT_Reclaim</code>
 <code>_EXFAT_Remove</code>

When accessing files via UserFS

- summarize full control flow of “open” syscall



Real world vulnerabilities - 1

- Vulnerable code ([CVE-2020-27904](#)), kernel FS

```

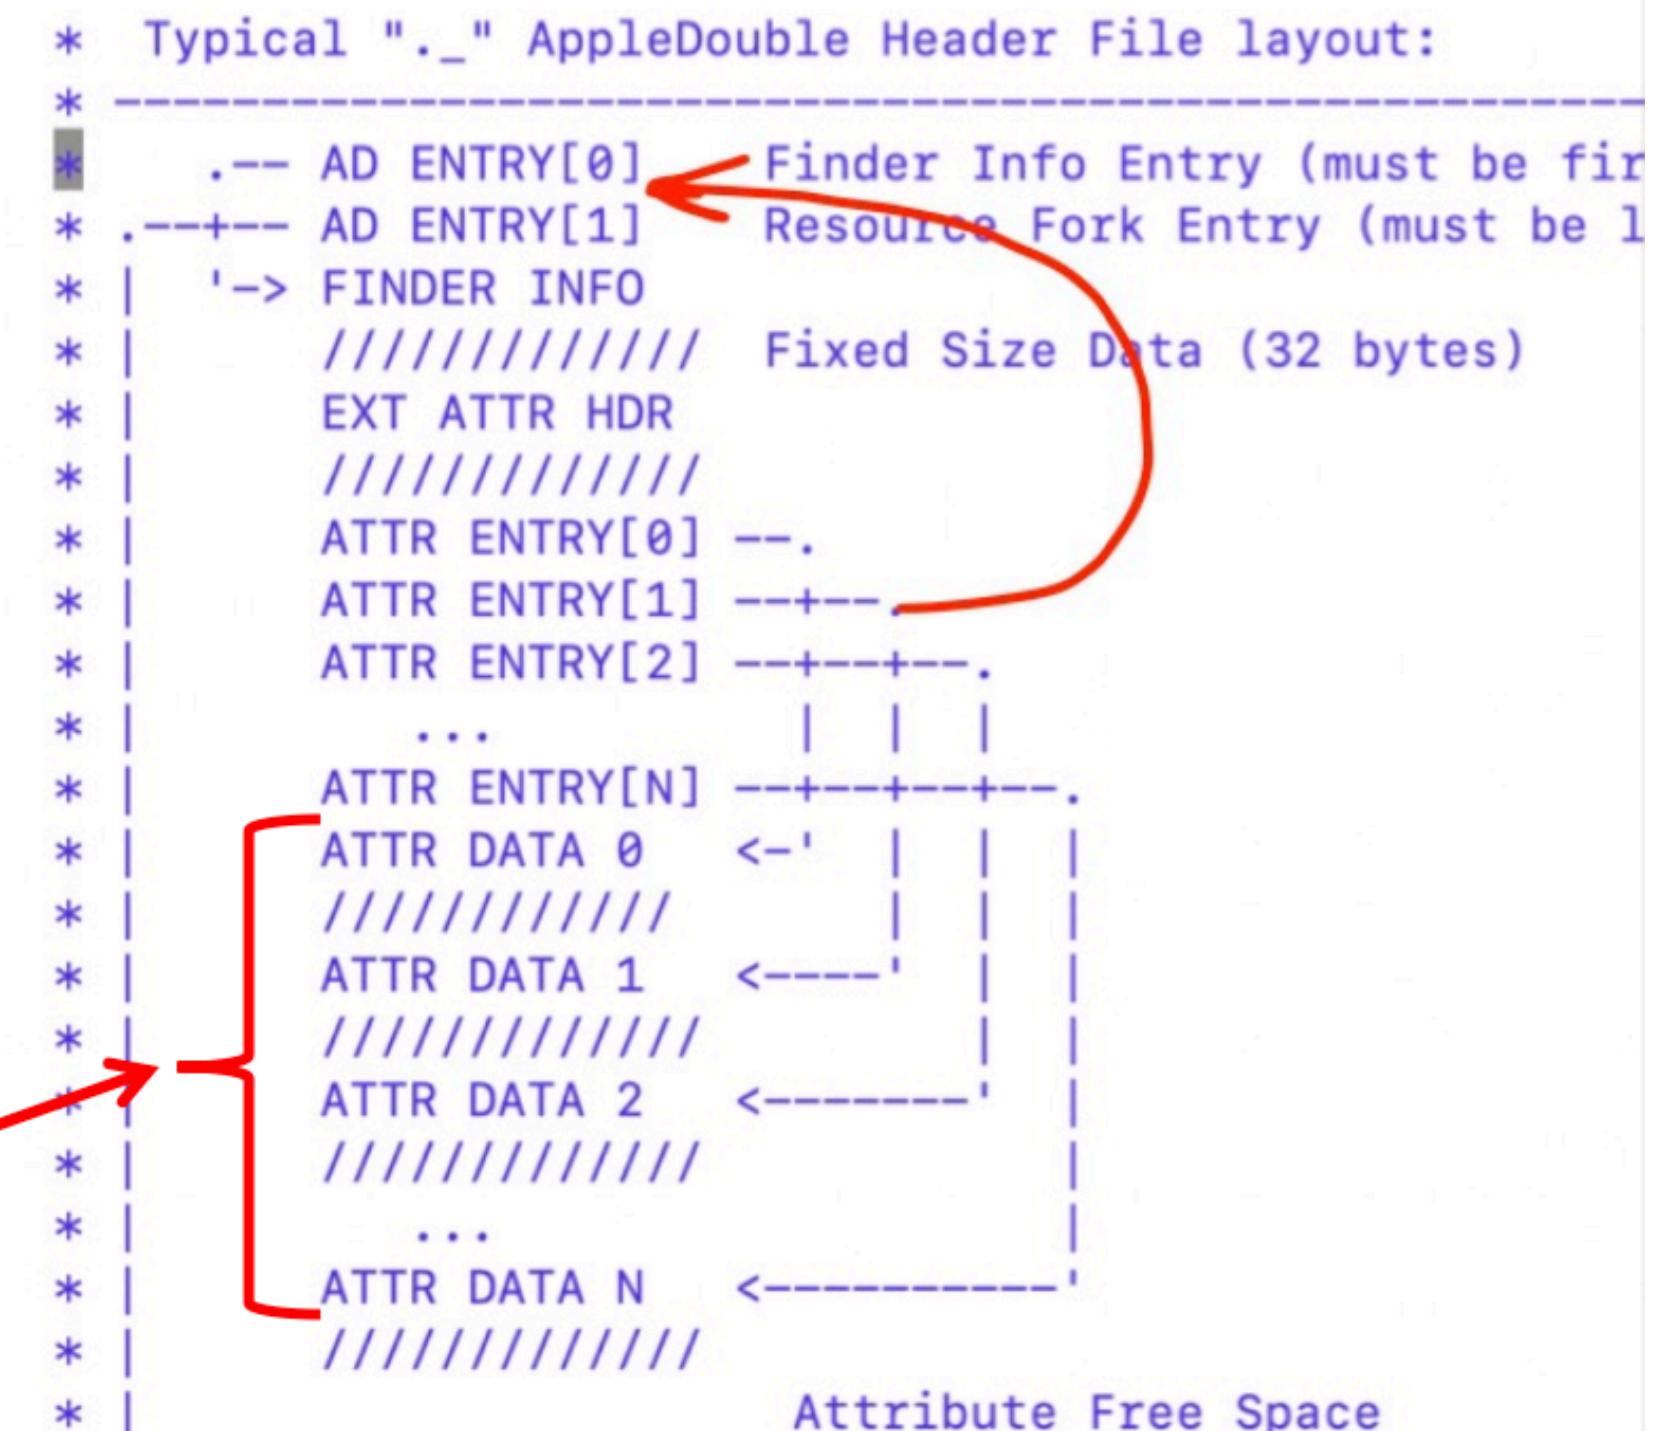
check_and_swap_attrhdr(attr_header_t *ah, attr_info_t *ainfop)
{
    /*
     * Make sure each of the attr_entry_t's fits within total_size.
     */
    buf_end = ainfop->rawdata + ah->total_size;
    count = ah->num_attrs;
    ae = (attr_entry_t *)(&ah[1]);

    for (i = 0; i < count; i++) {
        /* Make sure the fixed-size part of this attr_entry_t fits. */
        if ((u_int8_t *) &ae[1] > buf_end) {
            return EINVAL;
        }

        /* Make sure the attribute content fits. */
        end = ae->offset + ae->length;
        if (end < ae->offset || end > ah->total_size) {
            return EINVAL;
        }

        ae = ATTR_NEXT(ae);
    }
}

```



MUST: ah->data_start <= attr entry offset <= ah->total_size

Real world vulnerabilities - 1

- The vulnerability is about xattr (Extended File Attributes).
- exfat filesystem doesn't support xattr.
- XNU provides a compatible layer, the vulnerable code exists in compatible code.
- The logic is similar in UserFS, but the code is completely rewritten.

```
void -[liveFSVolume createAppleDoubleManagerIfNeeded](liveFSVolume *self, SEL a2)
{
    if ( (self->volCapInterfaces & 0x4000) != 0 ) {
        os_log_debug(OS_LOG_DEFAULT, "LFM: volume %@ supports xattrs natively");
    }
    else {
        v3 = +[LiveFSAppleDoubleManager AppleDoubleManagerForMount:self];
        self->appleDoubleManager = v3;
        os_log_debug(OS_LOG_DEFAULT, "LFM: created AppleDouble manager instance for volume %@", self);
        self->volCapInterfaces &= ~0x40000uLL;
    }
}
```

Real world vulnerabilities - 1

- UserFS version of **CVE-2020-27904**, i.e., CVE-2022-42861
- The rewritten code is LiveFS`-[LiveFSAppleDouble loadAttrHeader]

```
● 55     while ( buf_end >= (char *)&ae_ptr[1] )
● 56     {
● 57         namelen = ae_ptr->namelen;
● 58         if ( &ae_ptr->name[namelen] > (u_int8_t *)buf_end )
● 59             break;
● 60         v19 = strnlen((const char *)ae_ptr->name, namelen);
● 61         v20 = ae_ptr->namelen;
● 62         if ( v19 != v20 - 1 )
● 63             break;
● 64         ae_offset = bswap32(ae_ptr->offset);
● 65         ae_length = bswap32(ae_ptr->length);
● 66         ae_ptr->offset = ae_offset;
● 67         ae_ptr->length = ae_length;
● 68         ae_ptr->flags = bswap32(ae_ptr->flags) >> 16;
● 69         v11 = __CFADD__(ae_length, ae_offset);
● 70         ae_end = ae_length + ae_offset;
● 71         if ( v11 || ae_end > filehdr->total_size || ae_offset < filehdr->data_start )
● 72             break;
● 73         ae_ptr = (attr_entry *)((char *)ae_ptr + (((_WORD)v20 + 14) & 0x1FC));
● 74         if ( (unsigned __int16)++entry_idx >= (unsigned int)num_attrs )
● 75             goto LABEL_17;
● 76 }
```

Real world vulnerabilities - 1

- This kernel bug impacts iOS too, but only the UserFS code is affected.
- The UserFS one still exists on iOS < 16.2 & 15.7.2
- It didn't get fixed until I reported it again.
- root cause of CVE-2022-42861: the bug-fix of UserFS lost **synchronization** with kernel FS

Real world vulnerabilities - 2

- While studying UserFS, I found another xattr bug in kernel FS.

```
+++ b/bsd/vfs/vfs_xattr.c
@@ -2938,9 +2938,16 @@ get_xattrinfo(vnode_t xp, int setting, attr_info_t *ainfop, vfs_context_t conte
                     delta, context);
         writesize = sizeof(attr_header_t);
     } else {
-        /* Create a new, empty resource fork. */
+        /* We are in case where existing resource fork of length 0, try to create a new, empty r
         rsrcfork_header_t *rsrcforkhdr;

+        /* Do we have enough space in the header buffer for empty resource fork */
+        if (filehdr->entries[1].offset + delta + sizeof(rsrcfork_header_t) > ainfop->iostsize) {
+            /* we do not have space, bail for now */
+            error = ENOATTR;
+            goto bail;
+
+        }
+
         vnode_setsize(xp, filehdr->entries[1].offset + delta, 0, context);
```

- I think it is CVE-2022-42842.
- The ability of this bug is limited.

Then `init_empty_resource_fork` will initialize the memory block to some fix values. A partially **controlable oob-write** occurs. We can write these bytes beyond the 64KB buffer.

00000ee0	00 00 00 00 01 00 00 00	01 00 00 00 00 00 00 00
00000ef0	00 1e 54 68 69 73 20 72	65 73 6f 75 72 63 65 20	..This resource
00000f00	66 6f 72 6b 20 69 6e 74	65 6e 74 69 6f 6e 61 6c	fork intentional
00000f10	6c 79 20 6c 65 66 74 20	62 6c 61 6e 6b 20 20 20	ly left blank
00000f20	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00
*			
00000fe0	00 00 00 00 01 00 00 00	01 00 00 00 00 00 00 00
00000ff0	00 1e 00 00 00 00 00 00	00 00 00 1c 00 1e ff ff

For example, we can control that

```
*(uint8_t *) (buffer + 0x10000) = 0xff;
*(uint16_t *) (buffer + 0x10000) = 0xffff;
*(uint32_t *) (buffer + 0x10000) = 0xffff1e00;
*(uint64_t *) (buffer + 0x10000) = 0xffff1e001c000000;
```

- Tales from the iOS/macOS Kernel Trenches - [@jaakerblom](#)
- With the exploit primitive (kmsg type confusion), I can exploit the kernel FS xattr bug.

Exploitation: CVE-30937

- kmsg corruption can be done by turning a non-complex kmsg into a complex one (i.e setting bits 0x80000000 - MACH_MSGH_BITS_COMPLEX)
- In that case, the controllable inline data of a message without any real descriptors will be interpreted as if it was in fact real descriptors

Real world vulnerabilities - 2

- But that primitive was fixed on iOS 15.2, so I can't exploit this bug easily on macOS >=12.1,>=13.

```
[+] kapi_write64 -> 0x1234
[+] kernel base 0xfffffe0013a68000, kernel slide 0xca64000
[+] verify kernel header
_mh_execute_header
0000: feedfacf 0100000c 00000002 0000000c
[+] kernel_proc 0xfffffe0017d7c4c8, self_proc 0xfffffe1513a749f0
[+] now everything is OK
[+] spawn a root shell ->

root@secs-MacBook-Pro ~ # id
uid=0(root) gid=0(wheel) groups=0(wheel),1(daemon),2(kmem),3(sys),4(tty),5(operator),
(staff),29(certusers),61(localaccounts),80(admin),701(com.apple.sharepoint.group.tor),
204(_developer),250(_analyticsusers),395(com.apple.access_ftp),398(com.apple.s_ssh),
400(com.apple.access_remote_ae)
root@secs-MacBook-Pro ~ # sw_vers
ProductName:    macOS
ProductVersion: 12.0.1
BuildVersion:   21A559
root@secs-MacBook-Pro ~ #
```

Real world vulnerabilities - 2

- Of course, there is a corresponding UserFS version (no idea of its CVE no.)
- LiveFS`-[LiveFSAppleDouble loadADHeader]

```
 305     ``_oword lvar3[128] = v100;
 306     ``_oword lvar3[144] = v101;
 307     ``_oword lvar3[48] = (*_oword *)buf;
 308     ``_oword lvar3[64] = v96;
 309     ``_oword lvar3[80] = v97;
 310     ``_oword lvar3[96] = v98;
 311     ``_qword lvar3[160] = v102;
 312     block_v86.lvar2 = (_int64)capture_err;
 313     dispatch_sync(v68, &block_v86);
 314     objc_release(v68);
 315     objc_release(header_end);
 316     delta -= 0x11ELL;           // sizeof(rsrcfork_header_t) = 0x11e
 317     bzero(fileHeader_2->appledouble.pad, delta);
 318     -[LiveFSAppleDouble initEmptyResourceFork:](
 319         self,
 320         "initEmptyResourceFork:",
 321         (char *)self->_filehdr + self->_filehdr->entries[1].offset + delta);
 322         self->_filehdr->entries[1].length = 0x11E;
 323         v64 = 4096LL;
 324     }
 325 }
```

Possible to pwn kernel via UserFS?

- User space oob bugs: The kernel version of them have been proved to be exploitable. So, they are exploitable, in theory. I didn't try.
- Kernel space race condition bug: It can be converted to kernel UAF. Though it is very hard to write a workable exploit for it, it is exploitable in theory. I didn't try.
- Sandbox: App cannot access lifs.kext directly.
- Chain the user space UserFS oob bug with the kernel space lifs UAF?
 - I think there is a chance to attack kernel via UserFS, at least, in theory (again...).

The changes in filesystem security model

- call stack
 - kernel FS: SYS_open -> msdos.kext
 - UserFS: SYS_open -> lifs.kext -> userspace daemons -> exfat.dylib
- lifs.kext is simple, just forwards syscall to userspace daemons.
- main exploit target
 - kernel FS: kernel extension
 - UserFS: userspace daemons
- If successfully exploited
 - kernel FS: kernel read/write
 - UserFS: takeover a sandboxed userspace process
- UserFS will reduce the impact of FS vulnerabilities.

disadvantages of UserFS

- UserFS is an extra feature. It doesn't replace (all the) kernel FS.
- Attack surface = kernel FS + UserFS
 - iOS: only apfs & hfs are kept in kernel
 - macOS: all filesystem extensions are reserved in kernel
- Maintaining two code bases with identical functionality, but it's hard to keep bug fixes in sync between them.
- XPC everywhere in UserFS, so, performance?
 - Compared to CPU, USB disk is too slow. The performance is acceptable.

Conclusion

- More modules, more bugs.
 - UserFS is an additional feature, kernel FS is still here, thus increasing the attack surface of the filesystem.
- There is a chance to break iOS through UserFS.
 - But when accessing files stored on USB disks, only UserFS takes effect. In this case, UserFS can significantly reduce the impact of filesystem vulnerabilities.
- Overall speaking, I think UserFS is a successful effort by Apple.



Thanks!

pattern-f (@pattern_F_)